

ELOSALES Sales Coach

Standard Mode — the manager’s view

This revision turns the Standard **Sessions** and **Analytics** tabs into a manager’s window on the team: open a rep → see their last-2-days recordings and their six skills as letter grades, plus what they’re strong at and where they need coaching. The goal from your PDF: coach the strugglers, replicate the closers. **Expert mode is untouched.**

15

FILES BUILT

2

STANDARD SCREENS REVISED

0

EXPERT SCREENS CHANGED

856

TESTS GREEN

Governed build under ThinkerThinker.md & CLAUDE.md (AMD-006 four-layer); **A18** is the primary asset (surfacing rep data to a leader — the label is the defense). Screens shown are faithful representations of the revised Standard UI — not live captures. All work is **BUILT** and gate-verified (tsc · lint · production build · 856 tests); it has **not** yet rendered in a live session, so it is not marked TESTED. Migration **0187** is UNAPPLIED — and **may not be worth applying at all**: auditing this build found that nothing in the product can play a recording, so the retention it adds guards an asset nobody can hear (p.3). Decide that first. Five decisions are open, and they are mostly one question: **what does this feature give the rep?**

01 What changed — line by line

Each part of your PDF, what it now does in Standard, and its honest status.

PART	WHAT CHANGED	WHY IT MATTERS	STATUS
S1 Sessions tab	For a MANAGER, the tab is a team roster → open a rep → their recordings from the last 2 days, each linking to the call for playback. A rep still sees their own browse-only view (A10).	The manager can see what a rep is actually saying, to coach through data.	✓ Built
S1b Retention	Recordings auto-delete after 2 days unless SAVED (by the rep or a manager). The audio is purged; the transcript & scores are kept. (Migration 0187 + a purge cron.)	Keeps the useful call, clears the noise — your “delete after 2 days unless saved”.	✓ Built
A2 Analytics tab	For a MANAGER, a team roster → open a rep → their profile: the six skills as letter grades, plus Strengths and Growth areas. The old ELO/team-aggregate stays Expert-only. A rep still sees their own six /10 scores.	Transparency on who’s doing well and who’s struggling; data to teach better.	✓ Built
A2b Letter grades	The six scores (Tone, Speed, Talk/Listen, Questions, Objection, Closing) render as A+/A/A-...D. Each grade carries its /10 basis; the floor is “D / growth-area”, never an F.	A11: a grade that shows its count is a mirror, not a bare verdict. A18: no punitive F.	✓ Built
A2c Strengths / growth	“Doing well” and “struggling on” are derived from those same graded scores — high tiers become Strengths, low tiers become Growth areas.	Same data, coaching framing (A18) — one source, not a second opinion.	✓ Built
G Manager gate	Reading a named rep’s scores/recordings requires being a manager (Sales-Coach admin or CEO/COO/admin) in the SAME company. Enforced by a pure, unit-tested authz function.	Per-rep visibility crosses the old self-only \$A18 line — so the gate is tested, not trusted.	✓ Built
H Honest states	A rep with too few sessions shows “still accumulating”, never a fabricated grade. A failed read shows an error, never “no data”.	\$3.5/\$3.6 (no fabricated measure) and \$3.4 (an error is not an empty).	✓ Built
G2 Survives 0187 lag	Recordings degrade to the plain 2-day window when the 0187 column isn’t applied yet, instead of the tab going dark. Save hides itself until it can actually work; a save attempt returns an honest “not available yet”.	The Jul-3 Team-Chat outage repeated itself here — code assuming a migration is live. Caught in self-audit, fixed.	✓ Fixed
X Expert mode	Every change is a new file, an additive DB column, a backward-compatible route param, or code strictly inside the Standard branch. The git diff shows zero Expert render-path lines changed.	Your hard constraint: Expert stays exactly the same.	✓ Verified
?2 Rep-saves-own	The save API accepts the rep, but there is no rep-facing Save button yet. Your PDF says “saved by the manager or user” — the “user” UI is a scope decision.	A10 + spec: does a rep get their own recordings-with-Save surface?	🕒 Your call
?3 One profile vs two	Built as two views (Analytics = scores, Sessions = recordings). My earlier recommendation was ONE unified profile. A deviation I’m flagging, not hiding.	AMD-006 L3: from a rep’s low grade, can a manager reach that rep’s recordings in one place?	🕒 Your call

PART	WHAT CHANGED	WHY IT MATTERS	STATUS
75 The letter itself	Your PDF specifies letter grades, and I built them. But a letter is the most familiar <i>ranking</i> instrument there is — and the framework’s test on leader-facing labels fails anything that invites comparison “even slightly”. Option: pair each letter with its tier word (“D · growth area”).	Your spec vs. the framework — surfaced for your ruling, not resolved by me.	 Your call

Standard = the simplified experience (default for new reps). **Expert** = today’s full system, deliberately unchanged.


Honest status: every item is **BUILT and gate-verified** (tsc, lint, production build, 849 tests). None has rendered in a live session, so nothing is marked TESTED — apply **0187** and drive it on staging to confirm. An audit (against the framework) found and **fixed** two violations I introduced — a \$3.4 error-as-empty and an A19 stale comment — and flagged the two decisions above (?2, ?3) rather than resolving them silently.

02 The new system — Standard screens (manager view)

Faithful representations of the revised Standard UI (not live captures). A rep’s own self-view is unchanged (A10).


S1 Sessions — team roster

manager opens a rep’s recent recordings




S1 Rep → recordings

last 2 days · Save keeps one longer




A2 Analytics — team roster

manager opens a rep’s skill profile



A2 Rep → profile

grades, each showing the count behind it



Two of these screens changed *after* the first draft of this report, because auditing the build against the actual framework text (rather than my memory of it) found two real faults. **The rep now sees the same letter grade their manager reads** — before, the letter and the strengths/growth call were reads about the rep that the rep had no way to see, and you cannot challenge a verdict that is invisible to you. **The manager now sees the count under every grade** — before, the person with authority saw the letter with its evidence stripped out, and could only take the System’s word for it. A count is something you can coach on, and something a rep can argue with; a bare letter is neither. Sample data shown; real grades come from the rep’s scored calls, with a “still accumulating” state until there’s enough.

03 Framework spine, audit & what's yours to decide

Built to the spec, governed by the framework — and honest about what isn't finished.

Framework the build satisfies

- **A18** — surfacing rep data to a manager, the label is the defense: Strengths / Growth areas, grades as targets, no F, no leaderboard.
- **A10** — no shadow read: the manager sees the identical numbers the rep sees; same endpoint, same computation.
- **A11** — mirror not verdict: each letter grade travels with its /10; strengths/growth are derived from those counts.
- **\$3.5 / \$3.6** — no fabricated measure; a thin sample shows "still accumulating".
- **\$3.4** — a failed read shows an error, not zeros (fixed in audit).
- **AMD-006** — built foundation-up: data → APIs → composition → surface; Expert isolation proven by diff.

Honest status

- **BUILT + gate-verified**: tsc, lint, production build, 856 tests green.
- **NOT TESTED** live: no route has run against a real DB/auth here.
- **0187 UNAPPLIED** — but no longer fatal: Analytics never needed it, and recordings now fall back to the 2-day window until you apply it.
- **Purge cron dormant** until wired (CRON_SECRET + a vercel entry, after 0187).
- **Grade bands** (the /10→letter cutoffs) are a one-line tunable constant, flagged for you.

Audit result (against the framework, actual files re-read)

- **Fixed** — \$3.4: manager views rendered fetch failures as empty/wrong content → now an honest error state.
- **Fixed** — A19: a stale Standard-Analytics comment → rewritten to the role-conditional behavior.
- **Fixed** — **the rep couldn't see the grade you see about them**. Found by reading the no-shadow-read clause at its source: it requires the same data at the *same level of detail*, and the letter + strengths/growth call existed only on your screen. Now on theirs too.
- **Fixed** — **you were shown the grade without its evidence**. The rep's own card lists the countable behavior; your profile of them listed only the letter. That is backwards — it hands authority a verdict and removes the part you'd actually coach with. The counts now sit under every grade.
- **Fixed** — migration-coupling (the Jul-3 outage class, repeated): recordings assumed 0187 was live and would have been hard-down until you applied it. Now degrades to the 2-day window. The degrade-vs-fail-loud rule is pure and unit-pinned — writing those tests corrected it twice (a bare code check would have swallowed a *different* column going missing; the write path fails with a different code entirely).
- **Your decision ?2** — rep-facing recording-Save UI (spec "or user" + A10). Additive if you want it.
- **Your decision ?3** — unify the two rep views into one profile (my ⑧ recommendation) vs keep them split (closer to your two PDF screenshots). A real refactor either way; your call.
- **Your decision ?7 — decide this one first. Nothing in this product can play a recording**. The audio is uploaded and stored, and read by nothing that renders a player. So a manager clicking a recording gets a transcript; Save preserves a file nobody can hear; the 2-day purge deletes one nobody could have heard. The audio is real and piling up regardless — storage and privacy exposure with zero value drawn from it. Either build playback (~60 lines), or stop storing the audio (which deletes 0187, the purge and Save along with it), or ship as-is knowingly.
- **Your decision ?8** — who may *un-save*? Your PDF names who may save and is silent on this, and I decided it silently: today a rep can un-save what their manager saved, erasing who saved it, and the purge then deletes it. Harmless while nothing can play the audio — destructive the moment ?7 builds playback. If you build playback, rule on this first.
- **Your decisions ?5 + ?6 are one question** — what does this feature give the rep? Right now: letter grades on their own screen, and a notice their manager reads their recordings. All cost, no benefit. ?5 asks whether a letter-shaped verdict should exist at all; ?6 says a metric must come with an offered next step. Rule on the absence and both resolve.

Next: **rule on ?7 first** (it may delete 0187, the purge and Save entirely) → if playback, rule ?8 before building it → then the ?5+?6 absence, then ?2 and ?3 → then apply 0187 if it survives and drive Sessions & Analytics as a manager to turn BUILT into TESTED. The pre-0187 fallback is the least-tested code here and is exactly what a manager hits today — watch it first. Full record, with every finding and its evidence: *docs/feature-specs/ELOSALES-STANDARD-REVISION.md* (section 7). Founder queue: *docs/FOUNDER-ACTION-QUEUE.md*.