

What you need to do

Everything below is genuinely yours — a decision or an ops action. None is blocked on code I could write. Each item names what to do, my recommendation where I have one, and why it matters, so most can be settled with a yes/no.

Generated 2026-07-20 · after all migrations were applied to 0187

● Quick / low-effort
● A decision to make
● Needs your review
● Business call

1 One env var turns on features that are already built

The highest return for the least effort. These are built, tested, and — as of today — their database tables are all applied. They are inert only because a secret isn't set.

☐ Set `CRON_SECRET` in your Vercel environment QUICK

One variable activates **two** dormant background jobs at the next deploy: the **task-overflow sweep** (flags missed deadlines — the product was blind to them) and the **durability sweep** (\$3.5 measurement). Both are built, scheduled in `vercel.json`, and their tables are live. Nothing else to wire.

☐ Add the scheduled-report cron to `vercel.json` QUICK · NOW UNBLOCKED

Add `{ "path": "/api/finance/reports/deliver-cron", "schedule": "0 5 * * *" }`. This was *unsafe* before today — it needs views from migration 0172, which I applied today — so it is safe to add now. Uses the same `CRON_SECRET`.

2 ELOSALES Standard revision — decisions that gate shipping it

The manager-transparency build (from your PDF) is built but paused on these. Each carries my recommendation — approve, or override.

☐ ⑦ Can a manager actually play a recording? DECIDE

Today nothing in the product plays a call — a manager clicking one gets a transcript. Your PDF says "recordings ... delete after 2 days," which implies audio.

My recommendation: **build playback** (~60 lines; the signed-URL helper already exists). Your own words point at audio. Alternatives: drop the audio entirely (also valid), or ship transcript-only as-is.

☐ ② **Give reps a "Save recording" button of their own?** **DECIDE**

Your PDF says recordings are kept "unless saved by the manager **or user**." The manager save exists; the rep-facing one doesn't yet.

My recommendation: **build it** — "or user" is in your spec.

☐ ③ **One rep profile, or two?** **DECIDE**

Skill grades and recordings can live on one unified rep profile page, or split across two.

My recommendation: **unified** — one place the manager reads a rep.

☐ ⑤ **Show a letter grade, or the raw count?** **DECIDE**

Skills can display as A+/A/A- letters, or as the underlying count ("asked for the close in 2 of 9 calls").

My recommendation: **let the count be the label**. A count can't be argued with; a letter is a verdict a manager can misuse (\$A11/A18).

☐ ⑥ / ⑩ **Connect the insight to an action** **DECIDE**

A manager sees "this rep struggles at closing" and then... stalls — there's no next step. This is the workflow-continuity gap AMD-006 exists for.

My recommendation: **link each weak skill to the coaching move the product already generates** (the After-Pitch "next door"). Do *not* build a separate static tips list — it would duplicate what exists.

☐ ⑧ **Who may un-save a recording?** **DECIDE**

Your PDF names who may save; it's silent on un-saving. Right now a rep can release what their manager saved.

My recommendation: **rep always wins** on their own recording, with append-only attribution (who saved/un-saved is never overwritten). Note: enforcing this properly needs a DB trigger, not just app code.

☐ ⑨ **Reconcile two "team" views** **DECIDE**

C.A.R.E shows the team as an **aggregate only** (by design — no individual call-out). The Sales Coach shows **named per-rep grades**. Both are defensible; they should be a conscious choice, not an accident.

A consistency call for you — transparency-by-name (Sales Coach) vs protect-the-individual (C.A.R.E). Your PDF's stated goal is "transparency on who is doing well," which leans toward named.

☐ **Then: smoke-test it as a manager** **REVIEW**

Once the above are settled, open Sessions & Analytics logged in as a manager and confirm the roster, grades, counts, recordings, Save, and the 2-day purge behave. This is the build's "BUILT → verified by you" step — I can't see the rendered screens.

3 Finance — a couple of judgment calls

The finance system is ~96% built. What remains is decisions, not code.

☐ **Tax report: how to net credit notes** DECIDE

The tax report can overstate tax owed when a credit note applies. Full write-up in `docs/financial-system/TAX-CREDIT-NOTE-NETTING-DECISION.md`.

My recommendation: **proportional netting**. One read.

☐ **Does a credit note return goods to inventory?** BUSINESS CALL

Today a credit note reverses revenue but leaves goods expensed — correct for services, wrong for a returned physical product. Whether a credit implies a physical return is your business rule (a refund for a damaged item the customer keeps is not a return).

☐ **Three "build later" finance features** BUSINESS CALL

Scenario modelling (best after the cash forecast is in real use), **multi-entity consolidation** (only if you run more than one legal entity — do you?), and an **integration** (Stripe / Plaid / QuickBooks — which first, if any?). Also: `.xlsx` export needs a new dependency; CSV works today.

Resolved today: the recurring-bill "31st drifts to the 28th" decision is **done** — its recommended fix (an anchor-day column) was migration 0186, which I applied today. No action needed.

4 Security — two items I deliberately did not touch

Both have ready text; both were withheld from autonomous action on purpose because they need your eyes.

☐ **events INSERT-policy residual** REVIEW

A MED-severity fix (`coach.dissect_generated`) with ready SQL in `AUDIT-2026-07-09-brain-injection.md`. Held because it edits the single most critical RLS policy in the core chain — a core-policy change deserves your review before it ships.

☐ **C.A.R.E prompt-injection defense** REVIEW

The C.A.R.E persona has no prompt-injection guard; a warmth-preserving instruction is drafted in the findings doc. Held because the persona is carefully tuned and I can't verify warmth headless — add it, then smoke-test the tone.

5 Infrastructure — one thing to merge

☐ **Add a CI gate that runs every migration on each PR** REVIEW & MERGE

Today a migration with a bad column reference passes every CI check and only fails on apply — that's exactly how the 0175 bug I fixed today reached the repo. I designed the fix in `docs/proposals/ci-migration-gate.md` (uses the Supabase CLI, since 110 migrations need the `auth` schema a bare Postgres lacks). I did **not** wire it — it touches CI config and I can't run it headless.

My recommendation: **merge it non-blocking first**, watch it go green on real runs, then flip to blocking. Say the word and I'll open the PR.

How this maps to what I did today: you asked to automate migrations — that's done (`npm run db:check | db:dry | db:verify | db:apply`), the whole database is at 0187, and applying it surfaced + fixed one real bug (0175). Everything on this page is what's left *after* that, and every item here needs a decision or an action only you can take.

Fuller detail on any item lives in **docs/FOUNDER-ACTION-QUEUE.md**. This page is the short version.